



Introduction to Neo4j at Foo Café

About Me

- Information Developer
- IT Project Manager
- Oracle DBA

maria.scharin@neo4j.com



Neo4j

- Creators of Neo4j, the world's leading graph database
- Swedish founders
- 260 employees; main offices in
 - Malmö (Eng. HQ)
 - San Mateo (HQ)
 - London





PINK PROGRAMMING

Pink Programming is a non-profit association whose goal is to increase the number of women to program. We work for a programming world where women are as obvious as men. Programming joy should be available to everyone.

What will we do?

- Relational databases and Neo4j
- Introduction to the property graph model
- Introduction to Cypher
- Some Neo4j concepts for non-programmers
- *Hello World!*

Programming is like cooking



Cranberry Muffins

You will need:


275g self raising
flour
75g dark brown soft
sugar
50g caster sugar
150ml milk
150ml vegetable oil
1 egg
150g sweetened
dried cranberries
1tsp cinnamon
sugar for sprinkling

1. Preheat the oven to 180c/ gas mark 4.
2. Line a muffin tin.
3. In a mixing bowl combine the flour, sugars and cinnamon.
3. In a jug whisk the milk, oil and egg together.
4. Pour in the wet ingredients into the dry and quickly mic until just combined.
5. Stir in the dried cranberries.
6. Divide the mixture equally between the cases and top each with a sprinkling of sugar.
7. Bake for 20 minutes.

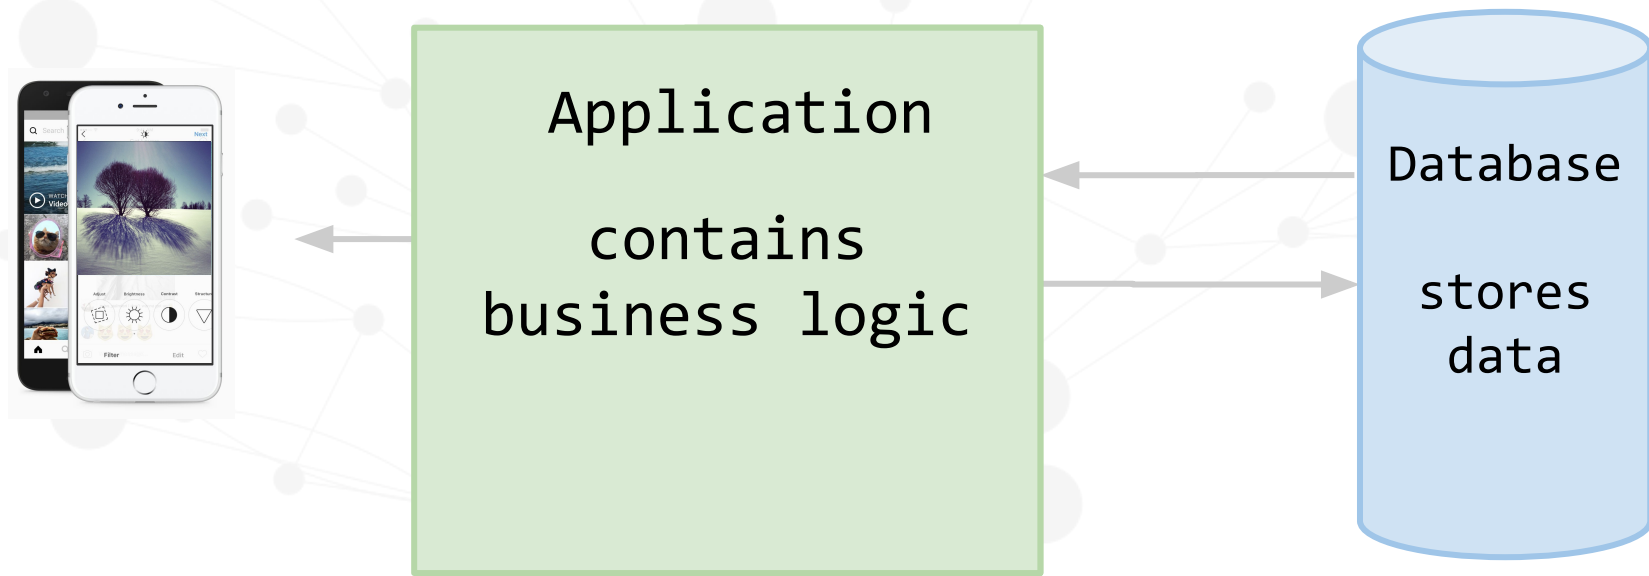
A computer program is like a recipe:

A set of instructions in a particular language that need to be followed in a particular order.

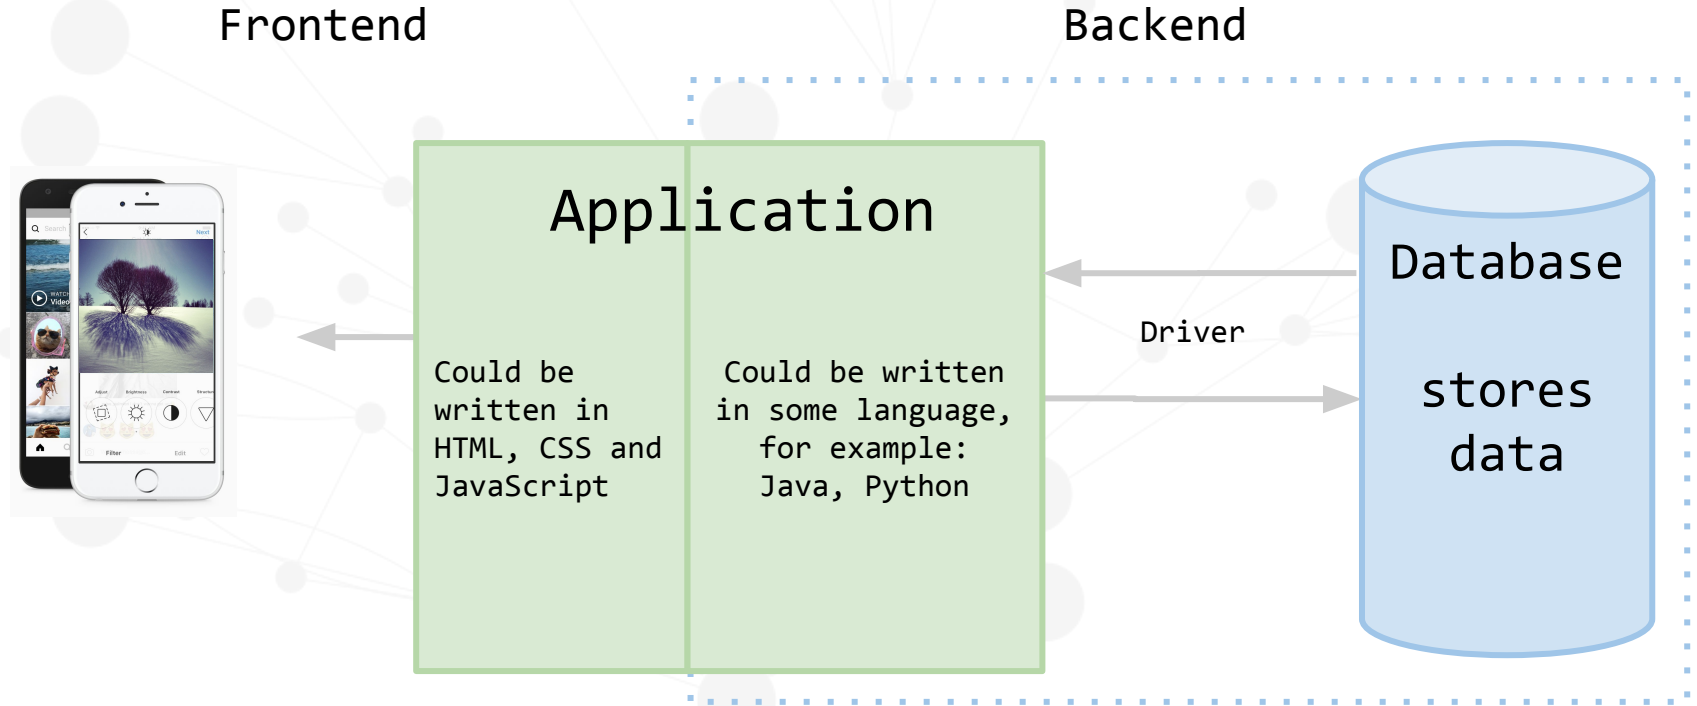
What's behind an app or a web site?



What's behind an app or a web site?



What's behind an app or a web site?



Languages



- Natural languages - Languages that we speak
- Programming languages
For example: Java, Python, JavaScript, C#
- Query languages - Used to “talk with” databases
For example: SQL, Cypher

Languages



Specific to programming languages

- There is no room for vagueness in programming. Either you do something or you don't!
- A programming language uses specific words to describe things and give instructions. We call those words *keywords*.
- Programming languages make things more efficient by using *variables*.
- Programming languages are very good at:
 - sorting large amounts of data
 - doing the same thing over and over (like a million times)
 - deciding whether to do something based on conditions (**if** somebody likes an instagram post **then** add a



Databases - Relational databases



Relational databases store data in tables.

Examples of relational databases:

- Oracle
- MSSQL
- MySql

FIRST_NAME	LAST_NAME	CITY
Maria	Scharin	Lund
Chelsi	Nolasko	Malmö
Myky	Tran	Lund
Emma	Herrlin	Malmö

You use SQL to write to and read from a relational database.

Databases - Relational databases



Benefits:

- They have been around for a long time; many people know them
- They store *structured data* very efficiently

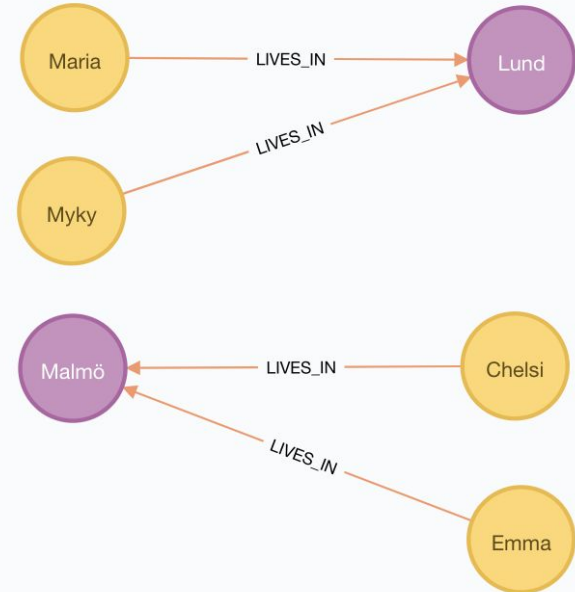
Drawbacks:

- They require specialized people, data modellers, to translate between the users and the developers. The model is abstracted.
- It is complicated to efficiently describe certain use cases; after the modelling is done it is hard for a non-IT person to recognize.

Databases - Neo4j



- Neo4j is a Graph Database.
- Neo4j stores data in nodes and relationships.
- You use Cypher to write to and read from Neo4j.

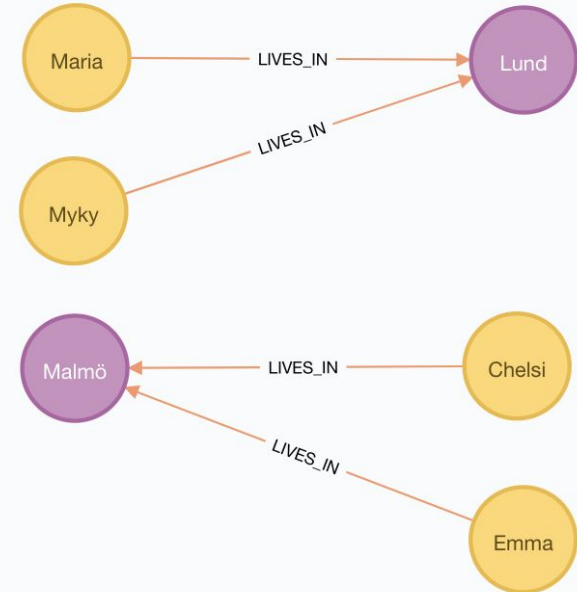


Databases - Neo4j



Benefits of graph databases:

- They describe things the way we see them; no abstraction needed.
- They store and retrieve connected data very efficiently.
- They are well suited for combining different types of data:
 - *Structured data* (structured in tables)
 - *Unstructured data* (emails, books etc.)
 - Data with various *data quality*.







Introduction to the property graph model

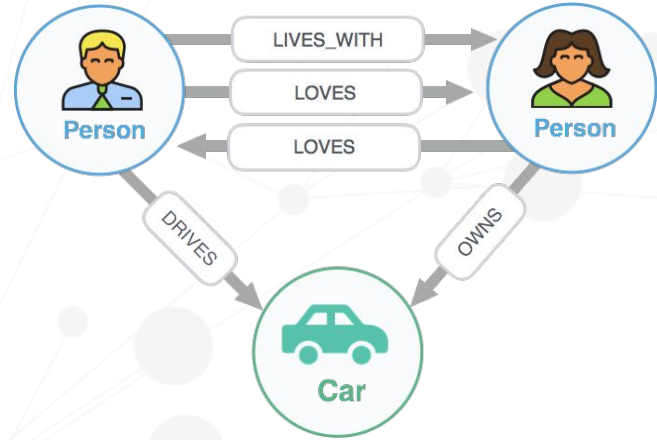
Nodes

- Nouns in your model
- Represent the objects or entities in the graph
- Can be *labeled*:
 - Person
 - Location
 - Residence
 - Business
- A node can have any number of labels



Relationships

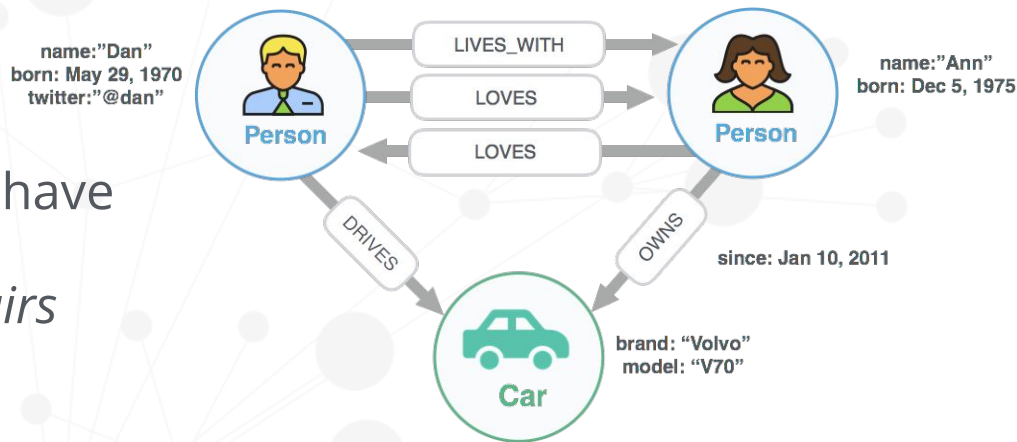
- Verbs in your model
- Represent the connection between nodes in the graph
- Has a type (exactly one):
 - LIVES_WITH
 - LOVES
 - OWNS
- Has a direction



Properties

- Used to *describe* nodes and relationships
- Nodes and relationships can have any number of properties
- Are described as *key/value pairs*

- name: "Dan"
- born: "May 29, 1970"
- since: "Jan 10, 2011"



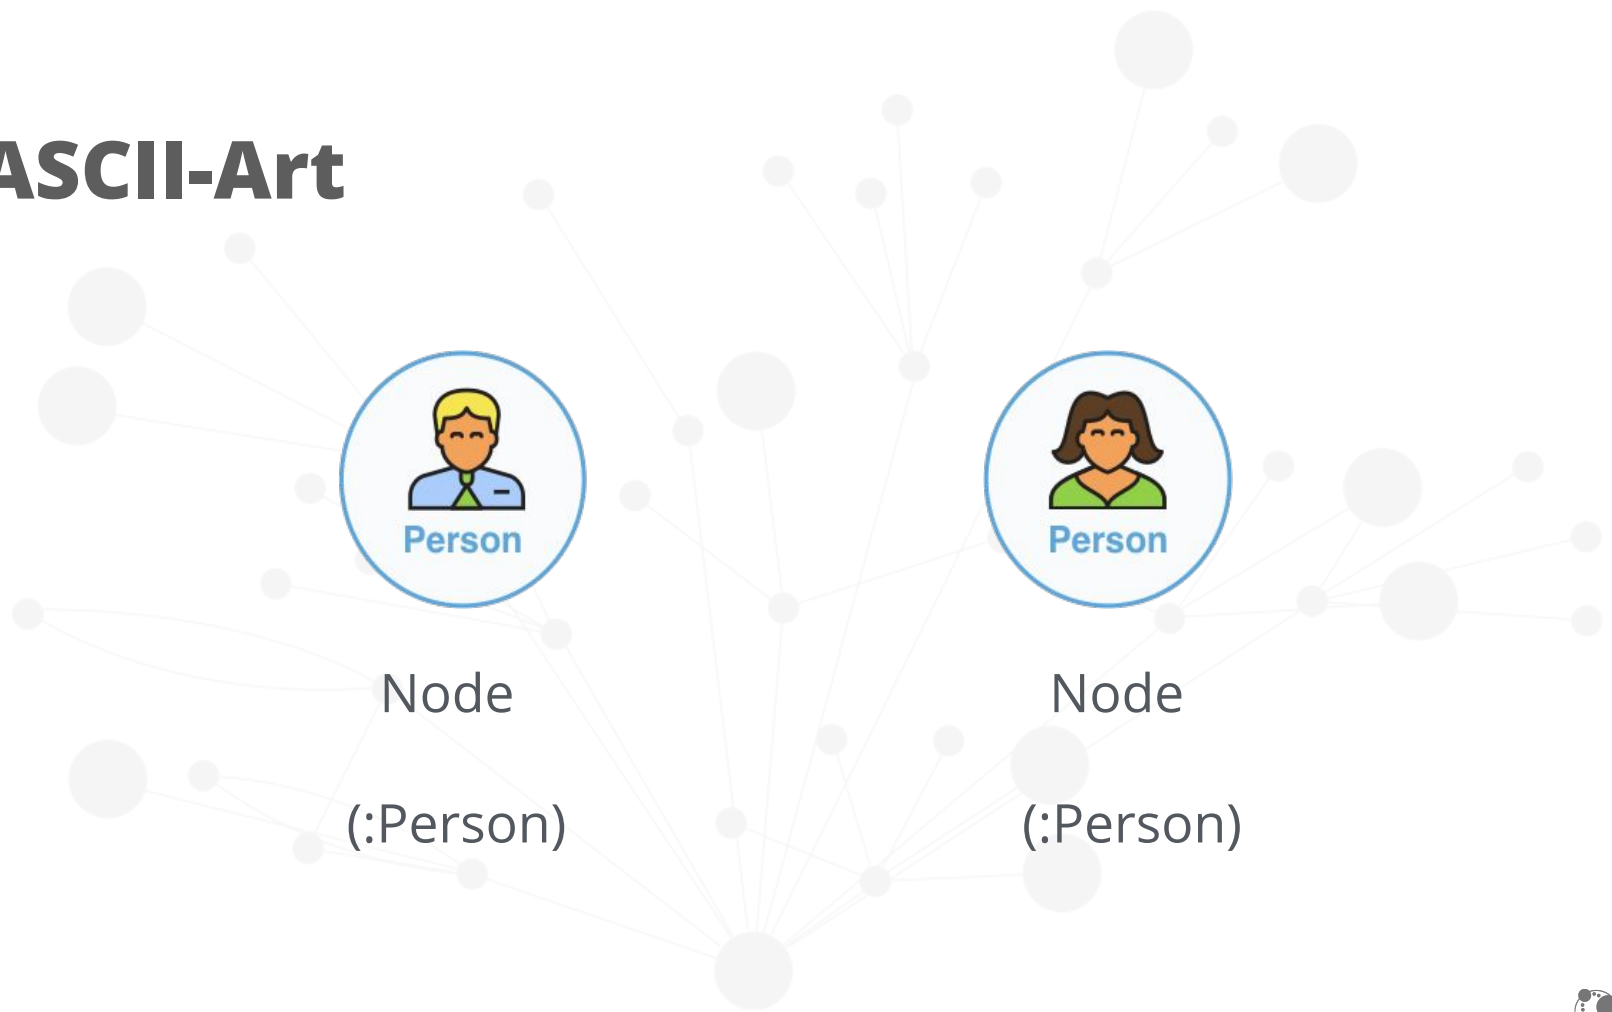
Introduction to Cypher

Cypher: The Graph Query Language

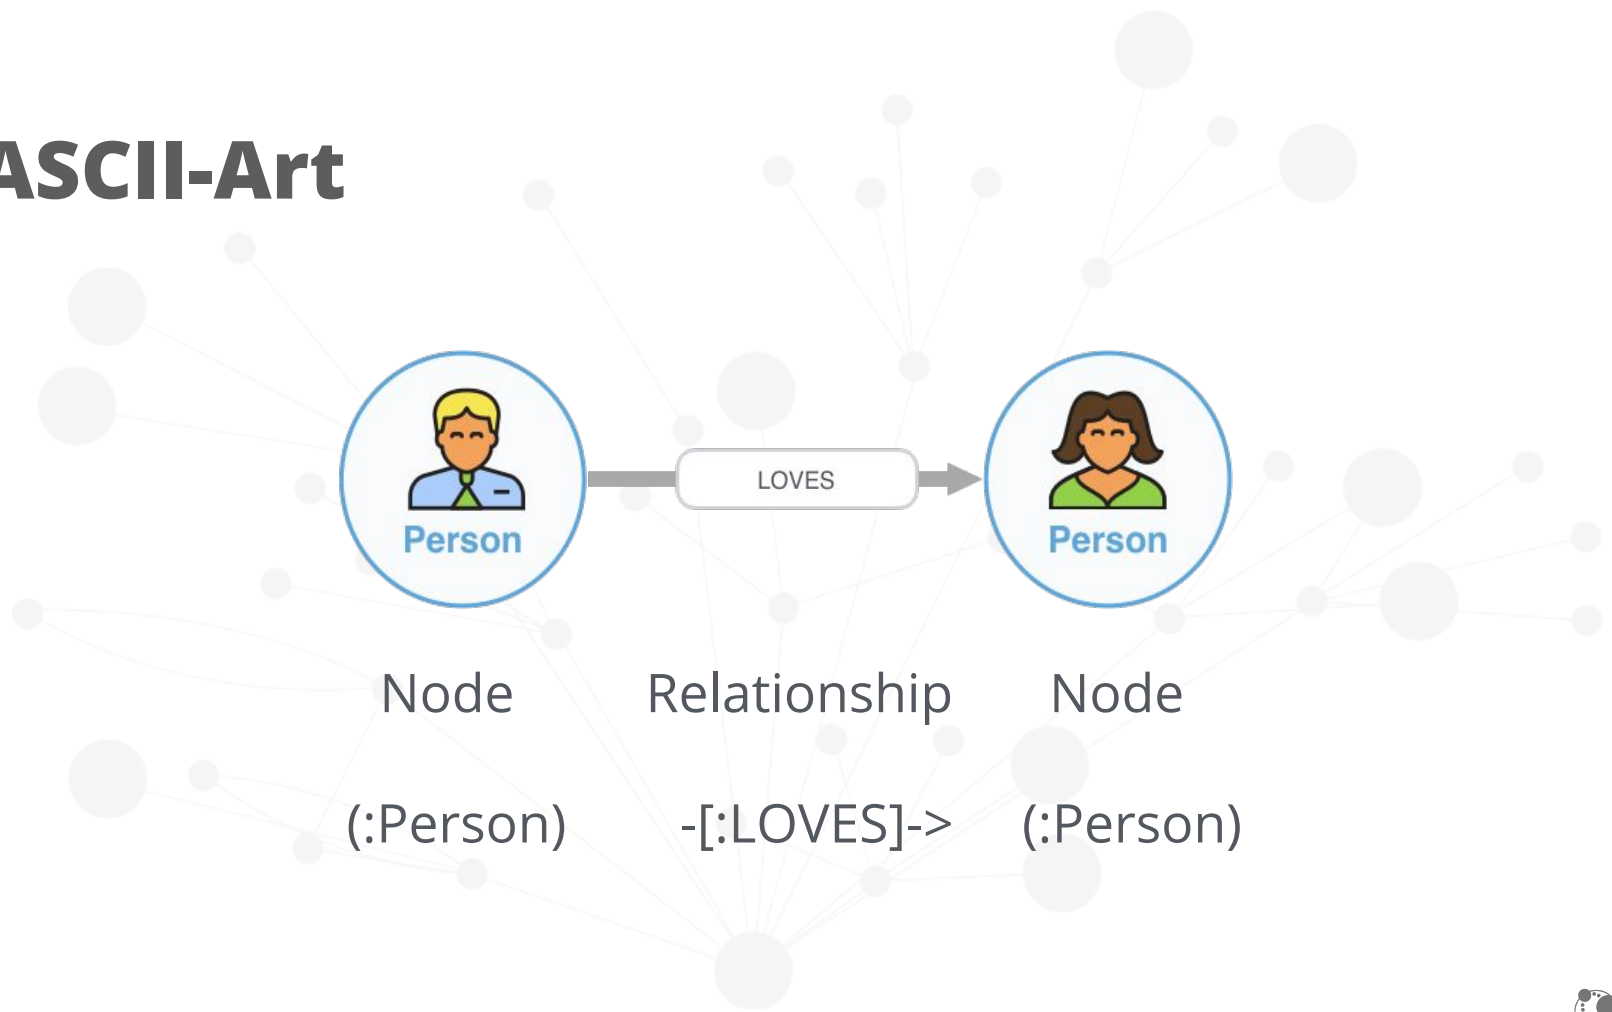
- Declarative - Focus on *what* to retrieve, not *how* to retrieve it
- Use pattern matching
- Intuitive and easy to learn (ASCII-Art)

~_(\ツ)_/~

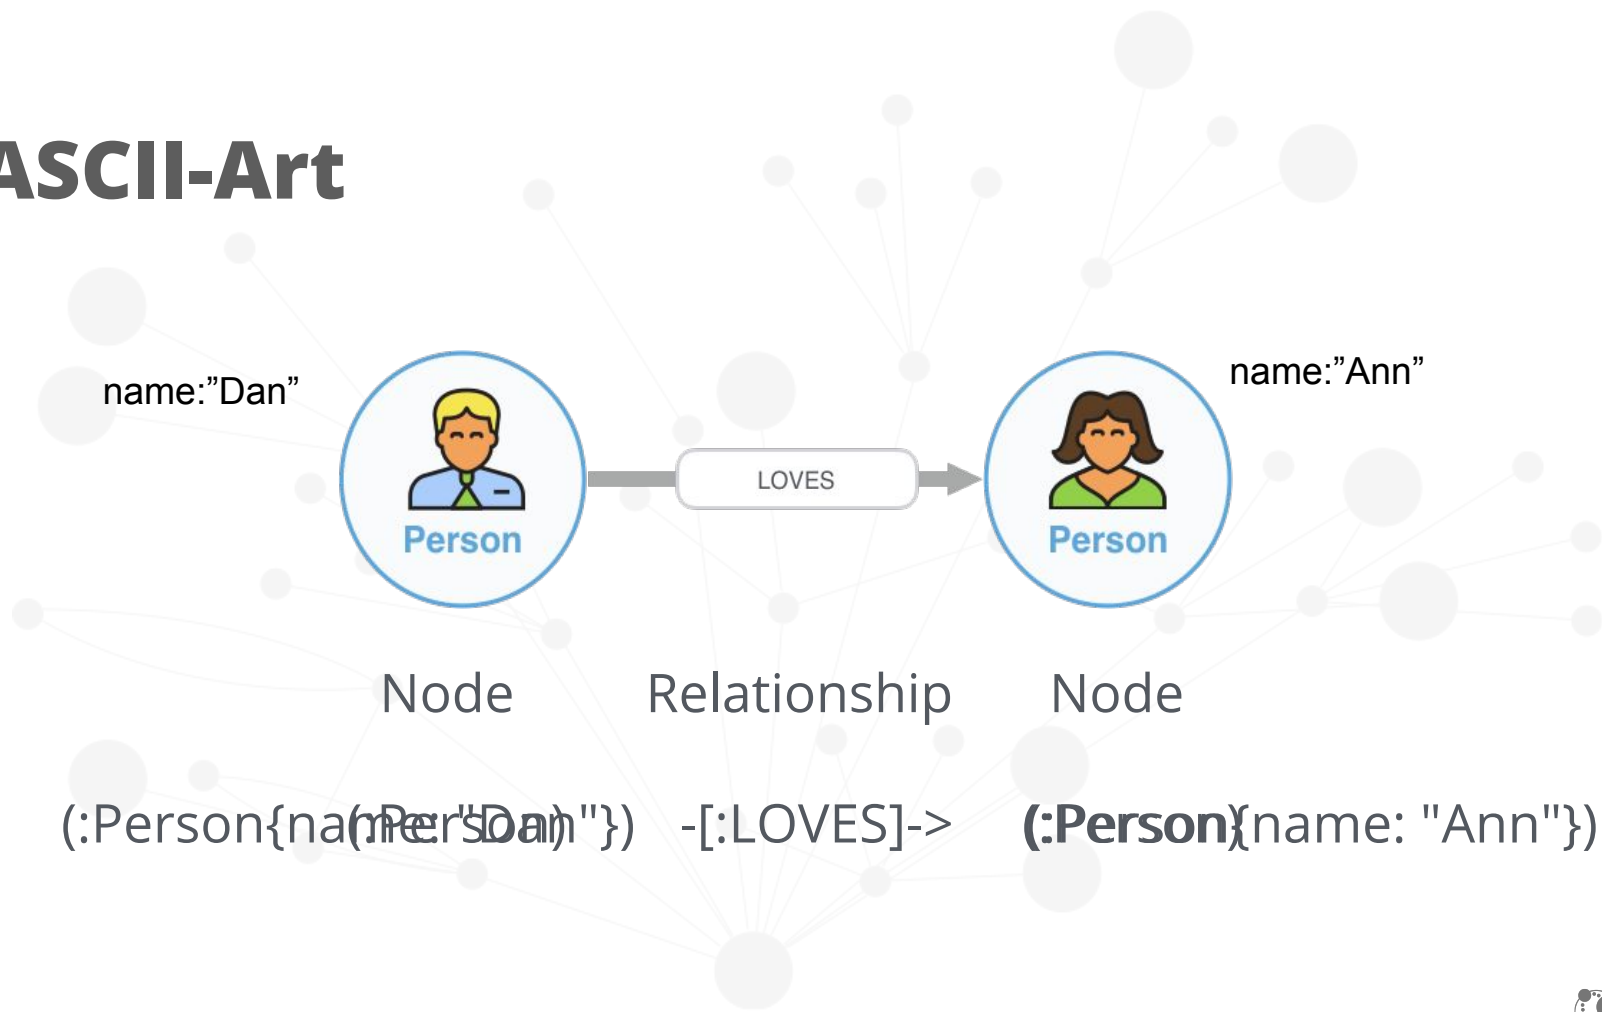
ASCII-Art



ASCII-Art




ASCII-Art



The CREATE keyword


name:"Dan"



Node

LOVES

Relationship



Node

name:"Ann"

```
CREATE (:Person{name: "Dan"}) -[:LOVES]-> (:Person{name: "Ann"})
```


CREATE nodes

CREATE ()

Create an "anonymous" node

CREATE (:Person)

Create a node with the label *Person*

CREATE (:Person{name: "Maria"})

Create a node with the label *Person* and a property called *name* with the value "Maria"

CREATE (:City{name: "Lund"})

Create a node with the label *City* and a property called *name* with the value "Lund"

CREATE relationships

Remember:

1. Every relationship must have a startnode and an end-node
2. Every relationship has a direction

CREATE ()-[]->()

Create an "anonymous" relationship between two "anonymous" nodes

CREATE ()-[:**LIVES_IN**]->()

Create an LIVES_IN relationship between two "anonymous" nodes

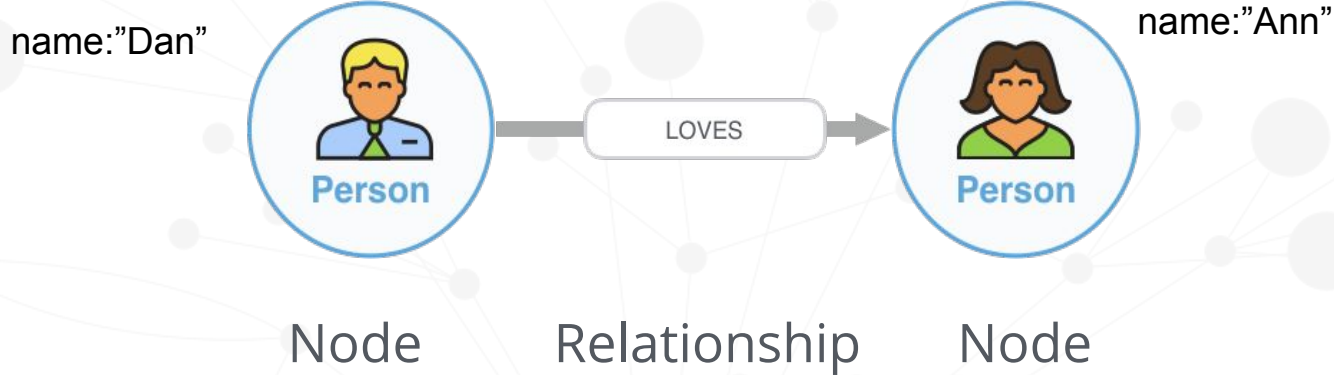
CREATE nodes and relationships

```
CREATE (:Person{name: "Maria"})-[:LIVES_IN]->(:City{name: "Lund"})
```

This does three things:

1. Create a node with the label Person and a property called name with the value "Nina"
2. Create a node with the label City and a property called name with the value "Lund"
3. Create a relationship of type LIVES_IN between Nina and Lund (with direction from Nina to Lund)

The MATCH and RETURN keywords



```
MATCH (p:Person{name: "Dan"}) -[:LOVES]-> (:Person{name: "Ann"})
```

```
RETURN p
```

We need to use a variable in order to be able to use what we have found later on. It can be a letter or a short word. In this example we chose "p" for "Person".

Using variables - 1

- A variable is a temporary container that we can put a value in, in order to use it later.
- In Cypher, we put values into variables in order to be able to use them later on in the query.
- We can give the variable any name we like, for example: `i`, `p`, `x`, `person`, `thisPerson`, `city`, `dan`....
- In these first examples, we use the variable to **RETURN** the value.

MATCH (p:Person) is the same as: **MATCH** (person:Person)
RETURN p **RETURN** person

Using variables - 2

```
MATCH (p:Person) -[:LOVES]-> (ann:Person{name:"Ann"})  
RETURN p, ann
```

Finds:

- all nodes with label *Person* (store these in the variable *p*)
- that has a *LOVES* relationship with
- another *Person* node that has a parameter *name* with a value "Ann" (store this in the variable *ann*)

Returns all the persons that we found (*p* and *ann*)

-- Everybody who loves a person called Ann!

Using variables - 3

```
MATCH (p:Person) -[:LIVES_IN]->(city:City)
RETURN p, city
```

Finds:

- all nodes with label *Person* (store these in the variable *p*)
- that has a *LIVES_IN* relationship with
- another node that has a label *City* (store these in the variable *c*)

Returns all the persons and cities that we found (*p* and *c*)

Using variables - 4

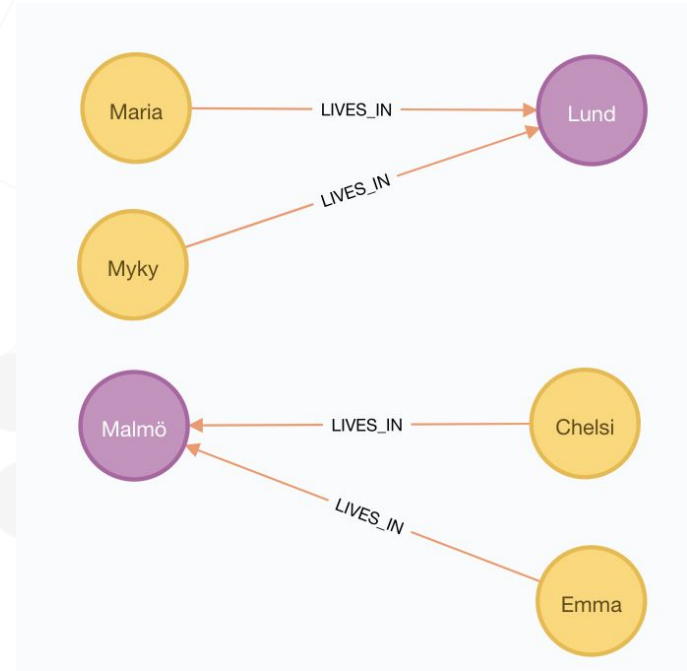
MATCH (:Person) - [**rel**:LIVES_IN] -> (:City)

RETURN **rel**

Finds:

- all nodes with label *Person*
- that has a *LIVES_IN* relationship with
- another node that has a label *City*

Returns all the relationships (**rel**)

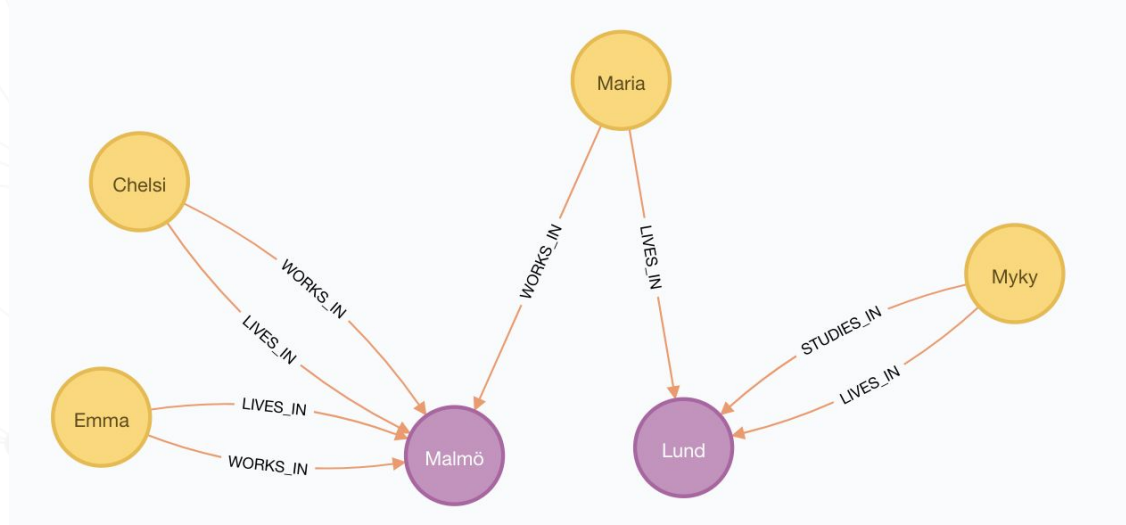


Using variables - 5

```
MATCH (:Person) - [rel:LIVES_IN] -> (:City)
```

```
RETURN rel
```

The query still only returns four relationships, since it specifically targets those with the type LIVES_IN



MATCH and CREATE in the same query

MATCH (p:Person{name:"Maria"})

MATCH (c:City{name:"Lund"})

CREATE (p)-[rel:LIVES_IN]->(c)

RETURN p, rel, c

What happens if
there are many
persons named
Maria or many cities
named Lund?

This does four things:

1. Find a node with label *Person* and name *Maria*; call it p
2. Find a node with label *City* and name *Lund*; call it c
3. Create a *LIVES_IN* relationship between p and c; call it rel
4. Return p, rel and c

RETURN values of properties on nodes

So far, we have returned the nodes, for example:

```
MATCH (p:Person)
```

```
RETURN p
```



If we instead want to get properties **in a list**, we use this syntax:

```
MATCH (p:Person)
```

```
RETURN p.name
```

p.name

"Maria"

"Myky"

RETURN values of properties on relationships

Find all *Person* nodes that has a *LIVES_IN* relationship with a *City* node and return the value of the properties *name* and *since*:

```
MATCH (p:Person) - [rel:LIVES_IN] - (c:City)
RETURN p.name, rel.since
```

p.name	rel.since
"Maria"	2004
"Myky"	2015

Watch out for case sensitivity!

Case sensitive:

Node labels

:Person is not the same as **:person**

Relationship types

:LIVES_IN is not the same as **:lives_in**

Property keys

Name is not the same as **name**

Watch out for case sensitivity!

Case insensitive: Cypher keywords

MATCH is the same as **MaTcH** but **MATCH** looks better

RETURN is the same as **return** but **RETURN** looks better

The background of the slide features a semi-transparent image of four people (three men and one woman) in a professional setting, looking at and working on laptops. Overlaid on this image is a network diagram consisting of numerous black circular nodes of varying sizes connected by thin, light-colored lines, creating a web-like structure across the right side of the slide.

Hello World!

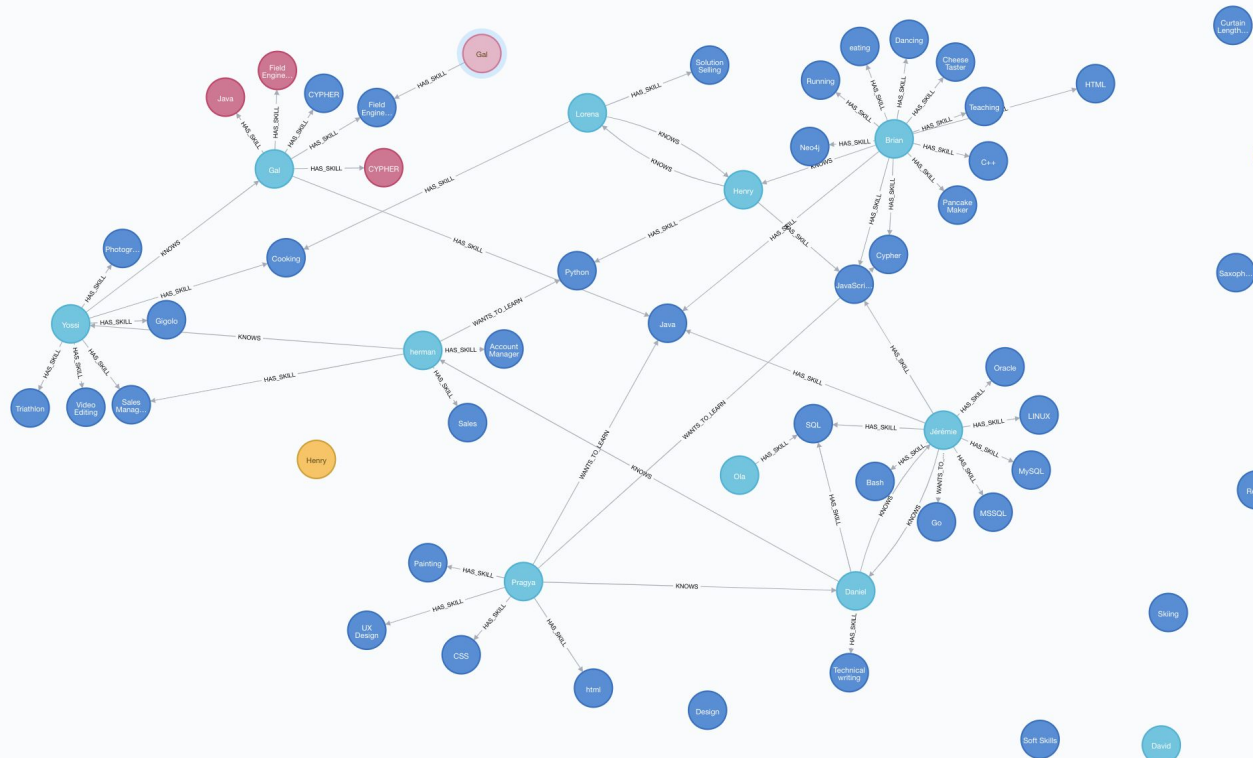
Let's graph our class!

Open document: <https://bit.ly/2GW4fbn>

Don't hesitate! This is just a play environment



Example from last class



Skill(81)

Person(4)

Person(18)

City(2)

SKILL(3)

person(3)

Let's try it - CREATE myself!

CREATE a person node within the graph to represent yourself!



Label of **:Person**

Property **name** should be **'your name'**

Welcome to the graph

CREATE statement:

```
CREATE (p:Person {name: 'Maria'})  
RETURN p
```

MATCH statement:

```
MATCH (p:Person {name: 'Maria'})  
RETURN p
```

Your turn: **CREATE** a node for yourself

CREATE yourself:

```
CREATE (p:Person {name: 'Your Name'})  
RETURN p
```



Your turn: **MATCH** and **RETURN** yourself

Find yourself:

```
MATCH (p:Person {name: 'Your Name'})
```

```
RETURN p;
```



SET and update properties

Find yourself and **SET** your city as a property.

```
MATCH (p:Person {name: 'Your Name'})  
SET p.city = 'Your City'  
RETURN p
```



Find your neighbour

Ask the person to your **left** their name.
Find the pair of you:

name: Your Name
city: Your City

Person

```
MATCH (p1:Person {name: 'Your Name'})
```

```
MATCH (p2:Person {name: 'Your Neighbour'})
```

```
RETURN p1, p2
```

Or in one statement:

```
MATCH (p1:Person {name: 'Your Name'}),  
        (p2:Person {name: 'Your Neighbour'})
```

```
RETURN p1, p2
```

name: Your Neighbour
city: Their City

Person


Creating relationships

Create a relationship between you and the person sitting next to you.

```
MATCH (p1:Person {name: 'Your Name'})
```

```
MATCH (p2:Person {name: 'Your Neighbour'})
```


```
CREATE (p1)-[:KNOWS]->(p2)
```



The Skills Graph

Finding skills

We've pre-populated the graph with a set of **Skills**.
See if you can write a query to find them.




Finding skills

Your query should look something like this:

```
MATCH (s:Skill)  
RETURN s
```

or

```
MATCH (s:Skill)  
RETURN s.name
```




s.name
"Java"
"Python"
"C#"
"Go"
"JavaScript"
"Scala"
"Bash"
"LINUX"
"Perl"
"HTML"
"CSS"
"Oracle"

Missing skills

The database contains a bunch of languages and databases.
Chances are you are missing your skills in there!
Can you create some?

NOTE: Each skill should have:

- a label called **Skill**
- a property called **name**



Missing skills

Your create statement could look something like this:

```
CREATE (:Skill { name: 'Your skill' })
```




Indexes and Constraints

You have a large amount of information of different kinds, for example in a book, and you need to find something very specific very quickly.

What will you do?



[illegible][illegible][illegible][illegible]

Constraints

Constraints are used to enforce business rules.

Examples:


- Each person must have a name
- Two persons cannot have the same ID number

Constraints

Remember: business logic is also enforced by the application code.

Database constraints are a way to make sure that the data follows the rules we have set up for the application.

We can also say that they *ensure the integrity of the data*.



Constraints

Constraints are used to enforce business rules.

Examples:

- Each person must have a name
- Two persons cannot have the same ID number

Creating a constraint

Let's create a constraint on `:Skill(name)`:

```
CREATE CONSTRAINT ON (s:Skill)  
ASSERT s.name IS UNIQUE
```



The MERGE clause

MERGE - Find or create

Let's try to create ourselves twice. What happens?

```
CREATE (s:Person {name: 'Your Name'})
```

MERGE - Find or create

```
MERGE (p:Person {name: 'Maria'})
```

```
RETURN p
```

This query does the following:

Search for a node with the label *Person* and a property called *name* with the value "Maria"

- If it finds it: **RETURN** it
- If it doesn't find it: **CREATE** it

MERGE - Find or create

```
MERGE (p:Person {name: 'Maria', nationality: "Swedish"})  
RETURN p
```

There is not a :Person node with name:'Maria' and nationality:"Swedish" in the graph, but there is a :Person node with name:'Maria'.

What do you think will happen here?

The MERGE Clause

```
MERGE (p:Person {name: 'Maria'})
```

```
SET p.nationality = "Swedish"
```

```
RETURN p
```

Merging relationships

To avoid **duplicate** relationships, use **MERGE**


MERGE (p1:Person {name: 'Your Name'})

MERGE (p2:Person {name: 'Your Neighbour'})

MERGE (p1)-[:KNOWS]->(p2)

Relationship is created uniquely by

- type
- direction (can be left off)
- properties



Your skills

Now let's associate you to some of those skills.

Can you write a query to add a **HAS_SKILL** relationship from you to whichever skills you have?

Hint: Don't forget to use the **MERGE** clause!

Your skills

Your query should look something like this:

```
MATCH (skill:Skill { name: 'Name of Skill' })
```

```
MATCH (me:Person { name: 'Your Name' })
```

```
MERGE (me)-[:HAS_SKILL]->(skill)
```

Multiple skills

What about if we want to do multiple skills in one query?

```
UNWIND ["skill 1", "skill 2", "skill 3"] as nameOfSkill
MATCH (skill:Skill { name: nameOfSkill })
MATCH (me:Person { name: 'Your Name' })
MERGE (me)-[:HAS_SKILL]->(skill)
```

Do your neighbours share any skills?

Now let's write a query to check if our neighbours have any of the same skills as us.

Do your neighbours share any skills?

Your query should look something like this:

```
MATCH (me:Person { name: 'Your Name' })-[:KNOWS]-(neighbour),  
        (neighbour)-[:HAS_SKILL]->(skill)<-[:HAS_SKILL]-(me)  
RETURN neighbour, skill
```


What skills do you want to learn?

Now let's associate you to some skills that you want to learn

Can you write a query to add a **WANTS_TO_LEARN** relationship from you to whichever skills you want to learn?

What skills do you want to learn?

Your query should look something like this:

```
MATCH (me:Person { name: 'Your Name' })  
MATCH (skill:Skill { name: 'Name of Skill' })  
MERGE (me)-[:WANTS_TO_LEARN]->(skill)  
RETURN me, skill
```

Can your friends teach you?

Can you write a query to find out if any of your friends already have the skills that you want to learn?

Can your friends teach you?

Your query should look something like this:

```
MATCH (me:Person { name: 'Your Name' })
```

```
MATCH
```

```
(friend:Person)-[:HAS_SKILL]-(skill)-[:WANTS_TO_LEARN]-(me)
```

```
RETURN friend, skill
```